FORMAL BUSINESS PLAN · REVISED INPUTS: OCTOBER 11, 2026

VendorLens

A cited, low-cost pre-contract vendor check.

AI vendor due diligence for procurement, finance, and marketplace trust-and-safety teams.

$1.008M	600	$250k
Base-case year-1 exit ARR	Base-case paying teams	Proposed pre-seed raise / grant

Prepared for the VendorLens hackathon submission. Financial figures are founder-supplied planning assumptions, not verified traction or historical results. Product-status statements reflect the founder's supplied description; no independent code audit was performed.

Prepared with cfo.ai · Exported October 12, 2026

1. Executive summary

VendorLens is an AI vendor due-diligence product that turns a company domain into a cited risk brief and a 0–100 score. It combines open-source intelligence (OSINT), technical domain checks, company enrichment, and web-search reports. Its assessment distinguishes two questions: **“Is this domain real and established?”** and **“Is this vendor’s security posture mature?”** It also provides a threat/reputation watch. These are risk indicators, not a certification of identity or security.

The target buyers are procurement teams, finance controllers, and marketplace trust-and-safety teams at small and mid-sized businesses onboarding 20–200 vendors per year. Starter costs $49 per month and Team costs $199 per month, both billed annually; Enterprise is custom-priced. The base case targets 600 paying teams at an average $140 per month by the end of year 1, equivalent to $84,000 MRR and $1.008 million annual recurring revenue (ARR).

The planned $250,000 pre-seed raise or grant funds approximately 15.6 months at a rounded $16,000 monthly cash-outflow budget, before revenue. The four specified monthly cost lines total $15,500. At the rounded budget, about 115 teams paying an average $140 per month produce enough revenue to cover that budget. This threshold assumes the specified costs remain adequate for actual usage; it is not a contribution-margin break-even calculation that independently scales investigation costs.

The shipped differentiation, as described by the founder, consists of evidence-linked briefs with confidence tags, a keyless low-cost OSINT stack, an open-source approach, and an on-device private browser verdict layer. **A cross-run cache is roadmap work. A retained vendor-risk graph is not shipped and is not claimed as an existing moat.**

2. Customer problem and market focus

Smaller organizations need a practical way to assess unfamiliar vendors before entering a commercial relationship or enabling a payee or seller. Public information is scattered across company records, domain infrastructure, security signals, and reputation reports. Manually assembling that evidence creates an inconsistent process and makes later review difficult.

VendorLens packages those checks into a brief an approver can read and trace back to public sources. The product is positioned as a **low-cost pre-contract check**, not a replacement for enterprise security ratings, formal know-your-business (KYB), document verification, or a security audit.

Buyer	Intended use	Plan fit (recommended)
Procurement teams	Document an initial vendor check before approval	Starter; Team for workflow integration
Finance controllers	Review unfamiliar counterparties and preserve cited evidence for payment approvals	Starter; Team for alerts or API use
Marketplace trust and safety	Bring domain and reputation signals into seller or partner onboarding	Team; Enterprise for governance requirements

The stated segment's 20–200 annual onboardings equate to approximately 1.7–16.7 new vendors per month. Starter’s 50 investigations therefore covers initial onboarding volume for many target accounts. Team conversion must be driven by API access and Slack alerts, not simply a larger quota. Any automated scheduled re-check workflow should be marketed only once implemented and validated.

No market-size estimate, measured willingness to pay, customer traction, acquisition cost, or retention history was supplied. The year-1 scenarios below are targets rather than evidence of demand.

3. Product and evidence workflow

A user enters a company domain. An agent runs the specified checks and assembles an evidence-based brief:

- **Domain identity and establishment:** RDAP/WHOIS, TLS, certificate-transparency history, and company enrichment.
- **Technical security posture:** DNS-over-HTTPS with SPF/DMARC, homepage security posture, and security.txt.
- **Threat and reputation signals:** urlscan, PhishTank reputation, and web-search reports.

The output includes a cited risk brief and a 0–100 score covering the two assessment layers, together with a threat/reputation watch. Every claim links to a public source and carries a confidence tag: **corroborated, single-source, or contradicted**. Lack of a public signal must not be presented as proof of safety; contradictory sources and inconclusive evidence should remain visible.

Recommended scoring policy: document whether a higher score means greater risk or greater trust, how the two layers combine, and how unavailable evidence affects the score. The supplied inputs do not define the score's direction, calibration, or weights; this plan does not invent them.

On-device verdict layer

The jev-web typed-decision model produces browser verdicts on-device: **no text leaves the device**. This privacy statement applies to the local verdict layer. It is not a claim that the agent's live OSINT and enrichment workflow makes no external requests; the live agent does use external data services.

Current behavior versus roadmap

Status	Capability or cost behavior
Current, founder-described	Cited briefs and confidence tags; low-cost OSINT checks; company enrichment on every live run; on-device private verdict layer
Roadmap	Gate Crustdata enrichment to Team+ to reduce Starter investigation cost; cross-run cache of past investigations
Not claimed as shipped	Retained cross-run vendor-risk graph of domains, registrants, or hosting

Pricing-tier entitlements are listed as supplied. The plan does not independently verify that every listed commercial feature is implemented or that Enterprise requirements have been satisfied.

4. Business model and pricing

Plan	Monthly equivalent	Investigations/ month	Entitlements
Starter	$49	50	Brief PDF export
Team	$199	500	API access, Slack alerts
Enterprise	Custom	Not specified	SSO, audit log, custom risk weights, DPA

All stated monthly prices use annual billing. Their annual equivalents are $588 for Starter and $2,388 for Team. Collection timing is not specified: annual billing supports upfront invoicing, but the cash plan must confirm actual payment terms before assuming collection at signup.

Assuming only Starter and Team, with no discounts or Enterprise revenue, the scenario average revenues imply Team shares of approximately 47.3% (conservative), 60.7% (base), and 67.3% (optimistic). These are mathematical plan-mix illustrations, not independently forecast customer preferences. Enterprise revenue, discounts, and churn would change the mix needed.

5. Unit economics and cost control

The founder assumes approximately $0.04 of blended model cost per investigation. Planner, observer, and analyst summary stages use cheap models via LiteLLM; most tools use free or keyless APIs. Crustdata company enrichment adds approximately $0.10 per investigation and **runs on every live investigation today**. The current blended model-plus-enrichment cost is therefore approximately $0.14 per live investigation.

Plan / behavior	Full-quota investigation cost	Revenue	Investigation-only margin
Starter, current	50 × $0.14 = $7	$49	85.7% (~86%)
Team, current	500 × $0.14 = $70	$199	64.8% (~65%)
Starter, roadmap enrichment gating	50 × $0.04 = $2	$49	95.9% (~96%)

These are the founder's quoted gross-margin assumptions, calculated on investigation costs only. They exclude hosting, queues, payment fees, support, and other service-delivery costs. Fully loaded accounting gross margin would be lower if those costs are classified as cost of revenue.

Usage-budget check: a $4,000 monthly inference-and-data budget at $0.14 per live investigation supports about 28,571 investigations, or 47.6 per team at 600 teams. This is an implied capacity calculation, not a stated usage forecast. At the base case’s illustrative Starter/Team mix, full quota usage would cost about $26,730 per month in model and enrichment fees alone—well above the $4,000 budget. The plan therefore requires materially lower average usage than full quotas, and that assumption must be tested rather than treated as proven.

Recommended controls: measure investigations and enrichment calls per account; track retries and cost per completed brief; monitor plan-level margins; and define overage or fair-use terms before heavy API usage emerges. Do not treat the proposed Team+ enrichment gate or cross-run cache as current savings.

6. Year-1 revenue scenarios

Case	Year-end paying teams	Average monthly revenue/team	Exit MRR	Exit ARR
Conservative	200	$120	$24,000	$288,000
Base	600	$140	$84,000	$1,008,000 (~$1.0M)
Optimistic	1,500	$150	$225,000	$2,700,000

ARR is the year-end recurring revenue run rate: teams × monthly average × 12. It is **not the amount of revenue earned or cash collected during year 1**. A monthly acquisition schedule, churn, renewals, and collection dates would be needed to calculate those figures. No unprovided straight-line customer ramp or year-end cash balance is presented as part of this plan.

Revenue averages are retained exactly as supplied. They should be validated through observed plan mix, discounts, and renewals. No numerical CAC, churn, lifetime value, or paid-marketing budget has been added.

7. Operating costs, capital, and break-even

Monthly base-case cost at month 12	Amount
Inference + data APIs	$4,000
Infrastructure: hosting and queues	$1,500
Founder salary	$8,000
Design contractor	$2,000
Sum of specified costs	**$15,500**
Rounded planning budget	~$16,000

The base case uses a lean founder-and-contractor staffing structure. No additional hires, acquisition spending, legal fees, employer taxes, benefits, insurance, payment charges, or contingency amounts are supplied. They are not assumed to be zero; they are unbudgeted and should be priced before using this plan as a cash commitment.

Funding and runway

The funding target is a $250,000 pre-seed raise or grant, not asserted secured capital. Assuming it is available as unrestricted operating cash, $250,000 ÷ $16,000 gives **15.625**

months (approximately 15.6) of zero-revenue runway. Using the exact $15,500 cost sum gives approximately 16.1 months. The headline keeps the supplied rounded $16,000 basis.

Here, “burn” is the expense/cash-outflow budget before customer receipts, not measured net cash burn. Revenue collections can extend runway; growth-driven costs and unbudgeted items can shorten it. This is not a forecast of the date of profitability.

Cost-budget coverage

$16,000 ÷ $140 = 114.3 teams, rounded up to **115 whole paying teams**. At 115 teams, average monthly revenue is $16,100. At a $120 average, the same calculation gives 134 teams (133.3 rounded up). The old ~450-team claim is removed.

This calculation already includes the $4,000 inference-and-data line in the $16,000 budget. It must not be combined with an additional deduction for the same investigation costs. A true contribution calculation instead starts with fixed costs and subtracts per-account variable costs from revenue; the usage forecast needed for that calculation was not supplied.

At the base month-12 target, $84,000 MRR less $15,500 specified monthly costs gives **$68,500** before unbudgeted expenses and taxes. This is a conditional arithmetic surplus, not forecast net profit or cash generation. The cost budget was supplied for the base case only and is not copied unchanged onto the conservative and optimistic cases.

8. Competitive landscape and positioning

Category / examples supplied	Primary emphasis	VendorLens position
Security ratings / TPRM: BitSight, SecurityScorecard	Enterprise security-risk and third-party risk programs	A lightweight pre-contract first check, not a replacement
Business-data providers: Crustdata, Clearbit	Company information and enrichment	An evidence brief combining company and domain signals; enrichment providers can be inputs
KYB: Veriff/Onfido KYB, Sumsub	Document-based identity checks	Public-source screening, not formal identity verification
Domain / URL tools: urlscan.io, VirusTotal, Google Safe Browsing	URL, page, or malicious-content signals	Combine signals into a vendor-oriented brief; such signals alone do not establish corporate identity

VendorLens differentiates through cited and auditable claims, inexpensive largely keyless checks, open-source delivery, and private on-device verdicts. This is a positioning comparison using founder-supplied competitor examples, not an independently verified feature or pricing audit.

The founder describes richer enterprise data options as often requiring annual contracts and costing dollars to tens of dollars per lookup. This varies by provider, product, and contract. Without vendor quotes, this plan does not assert a universal price, contract requirement, or per-lookup comparison. Likewise, the keyless OSINT stack does not mean that the paid Crustdata enrichment call is keyless or free.

9. Differentiation, defensibility, and roadmap

- **Cited evidence with confidence tags:** every claim links to a public source and distinguishes corroborated, single-source, and contradicted evidence, making the output reviewable.
- **Keyless, low-cost OSINT stack:** most checks use free or keyless APIs and low-cost models, enabling an accessible price without an enterprise data contract for the core OSINT checks. Paid enrichment remains part of current live-run cost.
- **Private on-device verdict layer:** verdict text does not leave the browser, providing a privacy-preserving local assessment option.

These are current differentiators, not proof of an uncopyable moat. Defensibility must be earned through evidence quality, workflow fit, reliability, and customer trust. No accumulated vendor-risk graph is asserted.

Roadmap only: a cross-run cache of past investigations, with appropriate freshness and privacy controls; and gating company enrichment to Team+ to reduce Starter COGS. A cache is not automatically a vendor-risk graph, and neither roadmap item is included in current cost savings.

10. Go-to-market and execution priorities

The following execution recommendations are proposals, not supplied commitments or shipped capabilities.

- **Validate the buyer and decision:** run pilots with procurement and finance teams around a specific pre-contract approval workflow. Test whether source citations change decisions and reduce manual research.
- **Lead with a sample brief:** demonstrate both assessment layers, confidence tags, and the source trail. Avoid selling the score as a guarantee.
- **Sell Team on workflow:** target API and Slack needs, particularly marketplace onboarding. Track upgrades to validate the base-case average revenue.

- **Use open-source distribution:** clear documentation and reproducible examples can support adoption and scrutiny. Establish support, licensing, and hosted-service boundaries before promising them.
- **Qualify Enterprise requirements:** confirm SSO, audit log, custom weights, and DPA requirements before committing to delivery or taking implementation fees.

Near-term execution should establish source reliability, scoring explanations, repeatable briefs, and observed cost per live investigation. The next commercial gate is pilot conversion and retention; the next financial gate is evidence that the $4,000 API budget can support the base case's actual usage. No dated delivery timetable is invented.

11. Risks and mitigations

Risk	Recommended response
Wrong or overstated conclusions	Preserve citations and confidence tags; separate evidence from inference; require human approval for consequential decisions
Incomplete or stale public information	Show collection time, missing evidence, and contradictions; do not equate “not found” with “safe”
API cost, rate limits, or provider failure	Track requests and retries; document limits; build graceful degradation; validate enrichment cost and availability
Low Team adoption or weak retention	Test API/Slack value and plan mix; measure repeat usage and renewal behavior rather than assuming large allowances create demand
Privacy, reputational, and legal exposure	Review source terms, data handling, score disclosures, correction processes, and contractual protections; scope on-device privacy precisely
Insufficient operating budget	Price omitted payroll burden, legal, support, payments, and acquisition costs before setting spend commitments

12. Milestones and management metrics

The supplied year-1 targets are 200, 600, or 1,500 paying teams under the three alternative cases. The budget-coverage milestone is approximately 115 teams at $140 average revenue, conditional on staying within the $16,000 cost budget.

Track monthly: paying accounts by plan; MRR and ARR; average revenue per account; customer additions, churn and renewals; Team upgrade rate; live investigations and enrichment calls per account; cost per completed brief; investigation-only and fully loaded

margins; API reliability; citation coverage and correction rate; cash, collections, and net cash burn.

Use those measurements to replace assumptions. A hackathon submission can demonstrate the workflow and cite the plan's commercial hypothesis; it should not present scenario targets as achieved traction.

13. Basis and limitations

This plan uses the founder's revised October 11, 2026 inputs for product scope, customer profile, pricing, per-investigation costs, revenue scenarios, operating costs, capital target, runway, and shipped-versus-roadmap positioning. Derived arithmetic is explained where used. No independent code review, competitor pricing verification, or market sizing was performed.

Unknowns include current traction, customer acquisition schedule, investigation utilization, churn, Enterprise pricing, payment terms, unbudgeted expenses, and funding receipt date. The document is a formal business-plan draft and submission artifact, not audited financial statements or a guarantee of commercial performance.