BUSINESS PLAN · OCTOBER 2026

VendorLens

AI vendor due diligence: a cited risk brief and 0–100 score from a single domain.

$1.0M	600	$16k	$250k
Base-case year-1 ARR	Base-case paying teams	Monthly burn at month 12	Pre-seed raise

Prepared with Ari, cfo.ai. All inputs are founder-supplied. Figures marked derived are calculated directly from those inputs. Figures marked illustrative rest on stated assumptions.

1. Executive summary

VendorLens gives small and mid-sized businesses the vendor due diligence that large companies pay risk teams and data vendors to do. A user enters a company domain. An AI agent runs open-source intelligence checks: company data, DNS, WHOIS, TLS, homepage analysis, and a search for scam reports. It then writes a **cited risk brief** with a **0–100 risk score**, and every claim in the brief links to its source. An on-device browser verdict layer, built on the jev-web typed-decision model, scores privately, so no data leaves the browser.

The product is for procurement teams, finance controllers, and marketplace trust-and-safety teams at SMBs that onboard 20–200 vendors a year. It is sold self-serve on annual plans: Starter at $49/mo, Team at $199/mo, and custom-priced Enterprise.

Scenario (year 1)	Paying teams	Avg $/team/mo	MRR (derived)	ARR
Conservative	200	$120	$24,000	$288k
Base	**600**	**$140**	**$84,000**	**~$1.0M**
Optimistic	1,500	$150	$225,000	$2.7M

The plan runs on a $250k pre-seed round and a lean cost base of about $16k/mo at month 12. The stated target is about 15 months to profitability at about 450 paying teams. Section 8 tests that target against the cost inputs, and the economics come out stronger than the headline suggests.

2. The problem

Every new vendor is a counterparty risk: invoice fraud, shell companies, lookalike domains, or sellers on a marketplace who disappear after payment. Enterprises handle this with third-party risk teams and paid data feeds. SMBs onboarding 20–200 vendors a year can't justify that overhead. In practice the check is a quick Google search, a glance at the website, and trust. That leaves three gaps:

- **Inconsistency:** each person checks different things, and the work is rarely documented.
- **No audit trail:** when a vendor goes bad, nobody can show what was checked or why the vendor was approved.
- **Unused technical signals:** domain age, WHOIS changes, certificate details, and DNS setup reveal a lot, but most finance and procurement staff can't read them.

3. The solution

Core product: the cited risk brief

One input, a company domain. One output, a brief built for a non-technical approver:

- **Company data:** identity, footprint, and consistency of public company information.
- **Infrastructure signals:** DNS, WHOIS (registration age and changes), TLS certificates.
- **Homepage analysis:** content quality, red-flag patterns, and whether claims match other sources.
- **Scam-report search:** public complaints and fraud reports.
- **0–100 risk score** with the drivers behind it, and **a citation on every claim**.

Private layer: on-device browser verdicts

A jev-web typed-decision model runs in the user's browser and gives verdicts with no data leaving the device. This answers the most common objection from privacy-sensitive buyers and lowers the cost of each verdict.

Workflow

Briefs export to PDF (Starter and up). Team adds API access, so checks can run inside onboarding and marketplace flows, plus Slack alerts. Enterprise adds SSO, audit logs, custom risk weights, and a DPA.

4. Customers

Segment	Job to be done	Likely plan
Procurement teams	Standard, documented check before a vendor is approved	Starter → Team
Finance controllers	Stop payment fraud before a new payee is set up; keep an audit trail	Starter
Marketplace trust & safety	Screen sellers and partners at onboarding, automatically	Team (API) → Enterprise

Ideal customer: an SMB doing 20–200 vendor onboardings a year, which works out to about 2–17 a month.

5. Pricing

Plan	Price (annual billing)	Investigations/ mo	Includes
Starter	$49/mo	50	Brief PDF export
Team	$199/mo	500	API access, Slack alerts
Enterprise	Custom	Custom	SSO, audit log, custom risk weights, DPA

Implied plan mix (derived). Treating revenue as Starter plus Team only, an average of $140/team needs about **39% Starter / 61% Team**. The conservative $120 average implies about 53/47, and the optimistic $150 about 33/67. Enterprise deals would lower the share of Team seats needed to hit each average.

Annual billing and cash. Annual plans collect a full year up front: $588 per Starter account and $2,388 per Team account. Cash then arrives well ahead of recognized revenue. This is a structural runway advantage and should stay the default.

6. Unit economics

Input	Value
Blended model cost per investigation	~$0.04
Crustdata enrichment per enriched investigation (Team+ only)	~$0.10
Gross margin at scale, Starter	~85%
Gross margin at scale, Team	~80%

Variable cost at full plan usage (derived)

Plan	Price	Max cost / mo	Margin on investigation cost only	Total cost to serve implied by stated margin
Starter	$49	50 × $0.04 = $2.00	95.9%	$7.35
Team	$199	500 × $0.14 = $70.00	64.8%	$39.80

Even a Starter customer using all 50 investigations costs about $2 a month in model fees. The 85% target leaves roughly $5 a month for infrastructure, payment processing, and

support. Team is different: a customer who uses all 500 enriched investigations costs $70, which brings margin down to about 65% before infrastructure. **The 80% Team margin holds as long as the average Team account uses no more than about 280 enriched investigations a month (about 57% of its allowance)**, or less once infrastructure is counted. Usage like that is very likely given the ICP (Section 9). Still, enrichment should run only when it's needed, and a fair-use policy should cap API-heavy marketplace accounts.

7. Cost structure

Monthly cost, base scenario at month 12	$/mo
Inference + data APIs	$4,000
Infrastructure	$1,500
Founder salary	$8,000
Contractor (design)	$2,000
Total burn	**~$15,500 (≈$16k)**

About two-thirds of the cost base is people ($10k), and the rest scales with usage. At base-case revenue of $84k MRR, inference, data, and infrastructure ($5.5k) come to about **6.5% of revenue**. That is well inside the 80–85% gross-margin targets.

8. Runway and path to profitability

15.6 mo	~115	~135
Runway at $16k/mo with zero revenue (derived)	Teams to cover $16k at $140 avg (derived)	Teams to cover $16k at $120 avg (derived)

The $250k pre-seed covers 15.6 months at $16k/mo even if VendorLens earns nothing. At the stated cost base, break-even takes about 115–135 paying teams. Each scenario's month-12 position:

Scenario	MRR at month 12	Burn	Monthly surplus (derived)
Conservative	$24,000	$15,500	+$8,500
Base	$84,000	$15,500	+$68,500
Optimistic	$225,000	$15,500*	+$209,500

*Usage costs rise with volume in the optimistic case. At the base rate of about $6.70 per team for inference and data, 1,500 teams adds about $6k/mo, which is immaterial to the result.

Illustrative year-1 cash path

Assumptions: paying teams grow in a straight line from 0 to each scenario's month-12 count. Fixed costs are $11.5k/mo (founder, contractor, infrastructure). Inference and data cost $6.67 per paying team per month. Revenue is counted as it's earned each month, which understates cash because annual plans prepay.

Scenario	Break-even month	Lowest cash balance	Cash at month 12
Conservative	Month 7	~$221k	~$259k
Base	Month 2	~$245k	~$632k
Optimistic	Month 1	$250k	~$1.5M

Reconciling the "15 months to profitability at ~450 teams" target. The inputs don't fully agree with each other:

- Fifteen months matches the zero-revenue runway ($250k ÷ $16k), not the time it takes to reach profitability.
- 450 teams at $140 is about $63k MRR, roughly 4× the stated $16k burn. Break-even at that burn is about 115–135 teams.
- The base case reaches 600 teams within 12 months, which already puts break-even before month 15.

For "profitable at 450 teams" to hold, monthly costs would need to be about $63k, meaning hiring that this plan doesn't include. One possible reading: **"self-funding by month 15 at ~450 teams, after reinvesting in a small team."** If that's the intent, the plan should list those hires. As written, the cost inputs say VendorLens can be **default-alive even in the conservative case**, and investors should hear that.

9. Market dynamics and what matters most

Usage vs. allowances. The ICP does 20–200 onboardings a year, or about 2–17 a month. Starter's 50 investigations already exceed what the largest ICP account needs for onboarding alone, and Team's 500 is about 30× it. Two consequences:

- **Good for margin:** real usage will sit far below allowances, so the 80–85% margins are conservative.
- **Risk to ARPU:** volume alone won't push anyone up to Team. The base case's $140 average (about 61% Team) rests on **API access and Slack alerts** being worth 4× the price. The strongest lever is **continuous monitoring of the existing vendor list** (re-checks and alerts on WHOIS, DNS, TLS, or reputation changes). That turns a one-time check into a recurring workflow and makes the allowances meaningful.

10. Moat

- **Cited-evidence briefs.** Every claim is sourced, so a brief can be defended in an audit, to a vendor, or in a dispute. Uncited AI summaries can't.
- **On-device private scoring.** The jev-web typed-decision verdict layer runs in the browser and no data leaves it. That's a privacy position cloud-only competitors can't easily copy, and it also costs nothing per verdict.
- **Accumulated vendor-risk graph.** Every investigation adds to a graph of domains, companies, infrastructure, and outcomes. Over time it catches repeat fraud patterns (shared registrants, hosting, templates) faster than one-off checks, and the advantage compounds with usage.

11. Go-to-market (recommended)

These are recommendations that fit the pricing and ICP above. They aren't part of the founder inputs.

- **Self-serve first.** Offer a free single-domain check that produces a shortened brief. Use the full cited brief and PDF export as the reasons to upgrade to Starter.
- **Finance and procurement content.** Publish vendor-fraud case breakdowns and onboarding checklists that use a VendorLens brief as the deliverable.
- **Outbound to marketplace trust-and-safety teams.** They are the natural API (Team) and Enterprise buyers, with the highest volume and biggest deals.
- **Integrations.** Slack now; procurement, AP, and onboarding tools next, so checks run automatically when a new vendor or payee is added.

12. Key risks and mitigations

Risk	Mitigation
Score errors (false positives harm legitimate vendors; false negatives harm customers)	Cite every claim, show score drivers, offer custom risk weights (Enterprise), and keep a human-in-the-loop disclaimer
AI-generated claims not supported by sources	Cite-or-drop rule: claims without a source are removed before the brief ships
Dependence on data providers (Crustdata)	Enrich only on Team+, only when needed; keep a second provider ready
Low ARPU if customers stay on Starter	Continuous monitoring, API, and Slack as upgrade drivers; track plan mix monthly
Established third-party risk and procurement suites moving down-market	Price, setup in minutes, cited briefs, on-device privacy

13. Milestones and metrics to track

Milestone	Why it matters
~115–135 paying teams	Break-even on the stated $16k cost base (derived)
200 paying teams	Conservative year-1 case: $288k ARR
Plan mix reaches about 60% Team	Needed for the $140 base-case average
600 paying teams	Base year-1 case: ~$1M ARR

Monthly dashboard: paying teams by plan, average revenue per team, investigations per account vs. allowance, enrichment rate on Team, gross margin by plan, logo churn, cash and months of runway.

14. Summary of inputs

Founder-supplied, used exactly as given: product scope; ICP (SMBs, 20–200 onboardings/yr); pricing (Starter $49 / 50 investigations, Team $199 / 500 investigations, Enterprise custom); model cost ~$0.04 per investigation; Crustdata ~$0.10 per enriched investigation (Team+); margins ~85% Starter / ~80% Team; scenarios 200×$120, 600×$140, 1,500×$150; month-12 costs $4k + $1.5k + $8k + $2k; $250k pre-seed; ~15 months to profitability at ~450 teams; three-part moat. Everything else is derived arithmetic or a labeled assumption or recommendation.